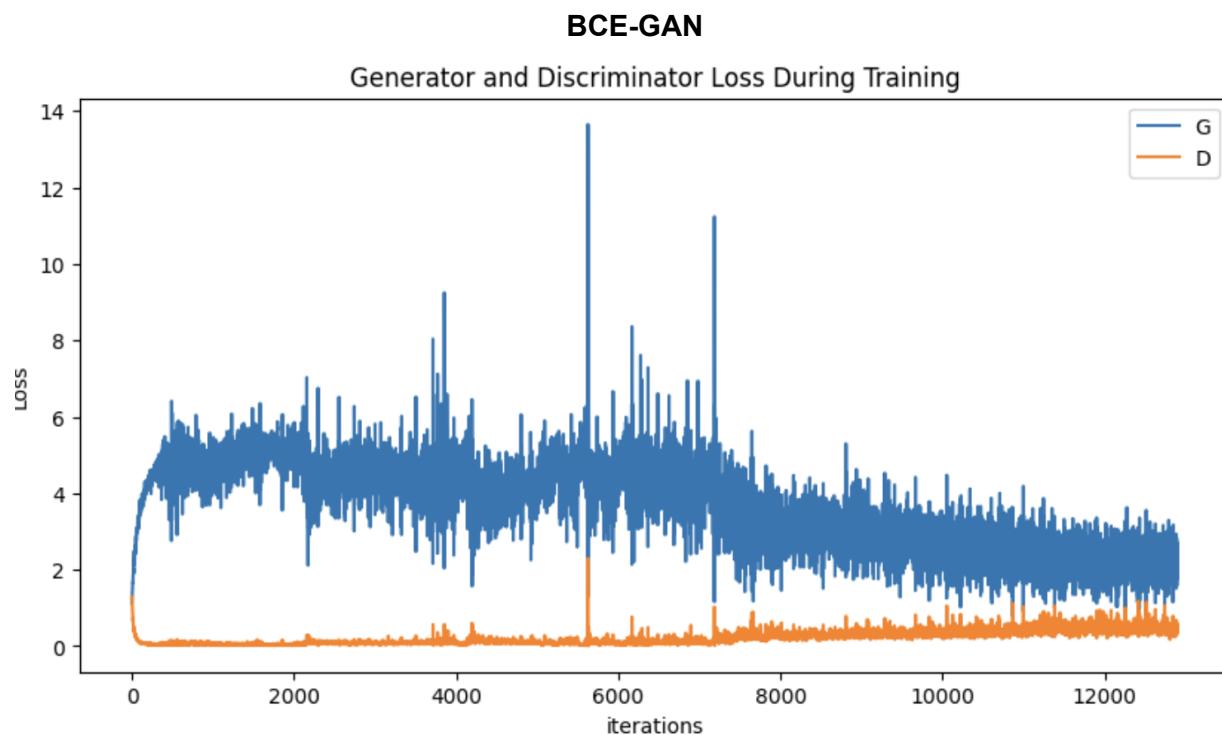


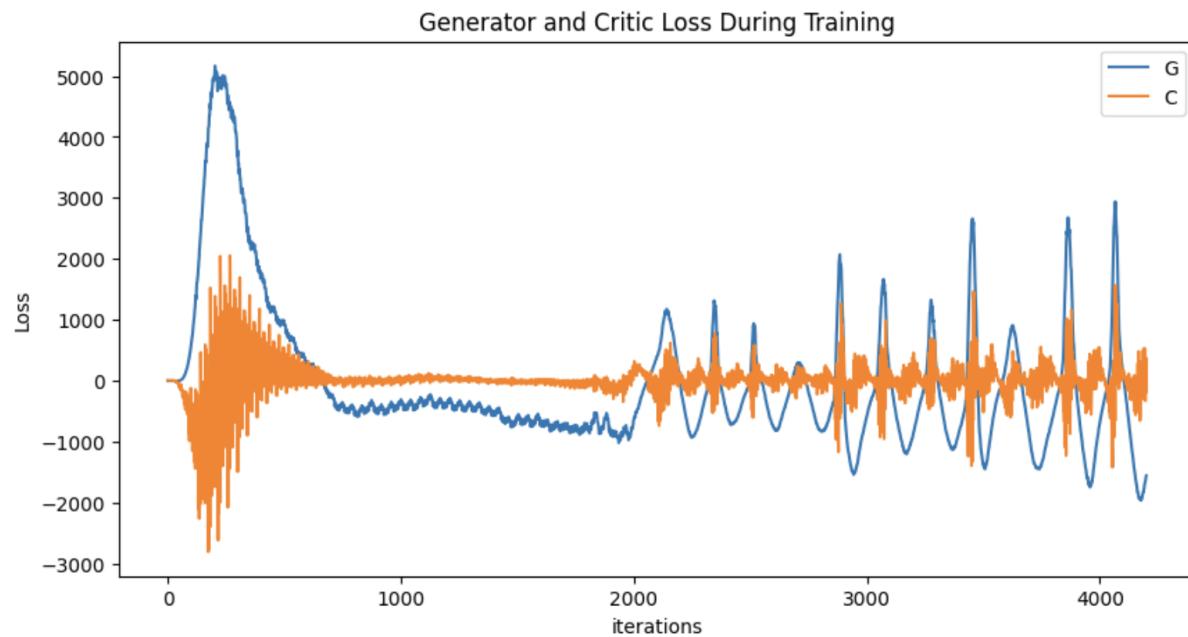
This assignment focuses on building two different GAN networks. The first is a BCE-GAN, which uses Binary Cross Entropy loss in the discriminator while the second is a W-GAN, which uses Wasserstein distance to evaluate loss (and gauge how far the generated distribution is from the training distribution). After training for 100 epochs on each network, the respective trained generators were used to generate 1,000 new “fake” images. Using these 1,000 images and the 1,000 evaluation images given to us, the FID scores were computed. Lastly, 16 fake images from each network are presented. At the end of the report, the code for this assignment is presented.

### Plot of Adversarial Losses over Training Iterations:



### W-GAN:

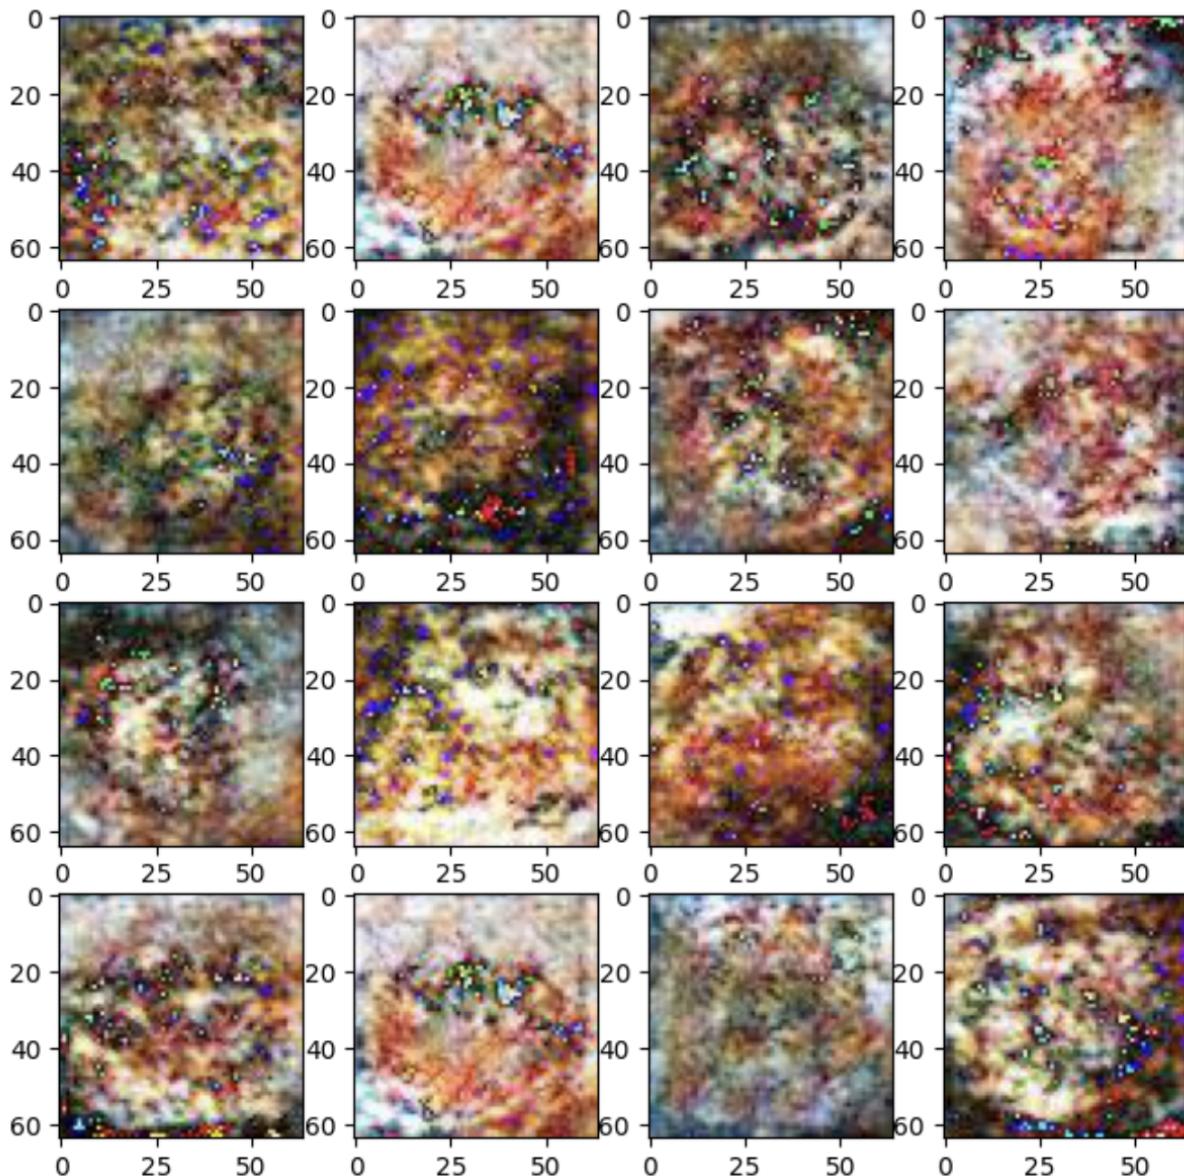
April 7, 2023



## Results:

**16 random images from the BCE-GAN Trained Generator:**

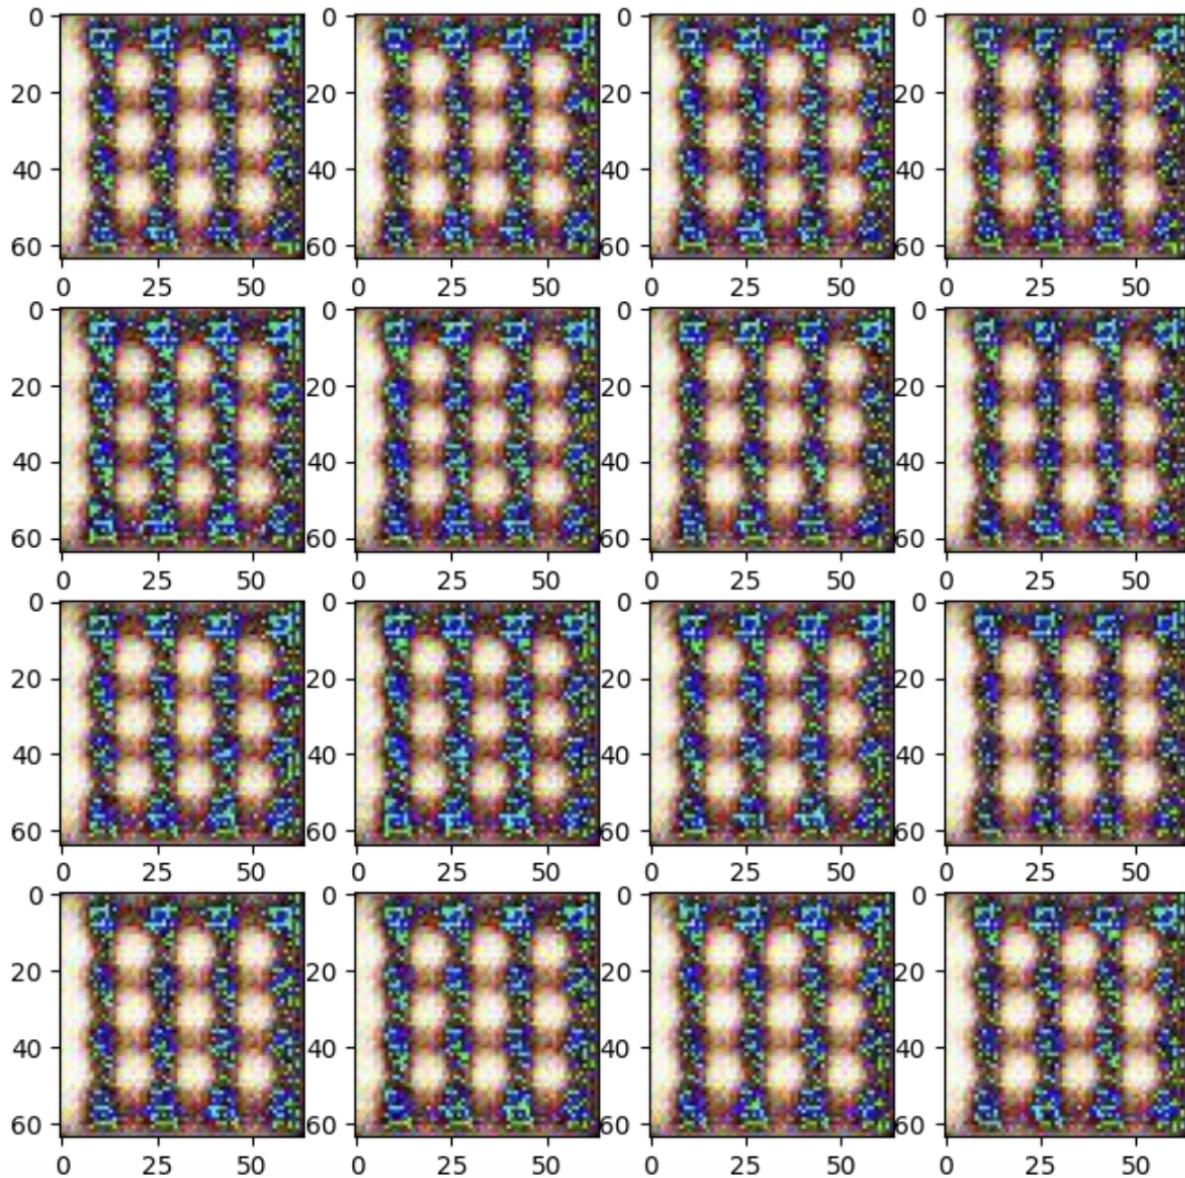
April 7, 2023

**FID for BCE-GAN Generator Images against 1000 Evaluation Images:**100% |  | 20/20

231.69390714090238

**16 random images from the W-GAN Trained Generator:**

April 7, 2023

**FID for W-GAN Generator Images against 1000 Evaluation Images:**

1000 | | 200

397.7164873260314

**Results Discussion:**

It is pretty apparent that the BCE-GAN is presenting better results. First of all, the generated images look almost like pizzas, but taken with a camera that has some distortion and noise. However, it is pretty apparent that the image quality is not as good as the original training set. Potentially, adding the skip connections to the BCE-GAN would improve the results. For W-GAN, it is apparent that the results suffer from mode collapse. However, it is interesting that the network starts to produce pizza-like colors (the yellowish/red circles) in the middle of the

April 7, 2023

image, while there is random noise in the background. Potentially, experimenting more with the gradient clipping penalty or lambda values can help improve the results for the W-GAN.

## Code:

### Dataloader

```

1 ### Create Dataloader
2 root = './'
3 folders = ['pizzas/train/0', 'pizzas/eval']
4 class MyDataset(torch.utils.data.Dataset):
5     def __init__(self, root, folder):
6         super(MyDataset).__init__()
7         self.path = root + folder
8         self.imgs = os.listdir(self.path)
9
10    for img in self.imgs:
11        if (img == ".DS_Store"):
12            self.imgs.remove(".DS_Store") #handle case when image isn't an image. Just remove it from the
13            #image list.
14            #print('removed DS store')
15    self.to_Tensor_and_Norm = tvt.Compose([tvt.ToTensor(), tvt.Normalize([0], [1])])
16    #print(len(self.imgs))
17
18    def __len__(self):
19        #return 102
20        return len(self.imgs)
21
22    def __getitem__(self, index):
23        PIL_img = Image.open(self.path + '/' + self.imgs[index])
24        torch_img = self.to_Tensor_and_Norm(PIL_img)
25        return torch_img
26
27 train_dataset = MyDataset(root, folders[0])
28 val_dataset = MyDataset(root, folders[1])
29
30 index = 2
31 print(train_dataset[index].shape)
32
33
34 train_dataloader = torch.utils.data.DataLoader(train_dataset, batch_size=64, num_workers = 0, drop_last=False)
35 val_dataloader = torch.utils.data.DataLoader(val_dataset, batch_size = 64, num_workers = 0, drop_last = False)
36

```

*Generator and Discriminator Network for BCE-GAN. Note that the code was slightly modified from Professor Kak's DLStudio code, but mostly the same.*

April 7, 2023

```

4 class DiscriminatorDG1(nn.Module):
5     """
6     This is an implementation of the DCGAN Discriminator. I refer to the DCGAN network topology as
7     the 4-2-1 network. Each layer of the Discriminator network carries out a strided
8     convolution with a 4x4 kernel, a 2x2 stride and a 1x1 padding for all but the final
9     layer. The output of the final convolutional layer is pushed through a sigmoid to yield
10    a scalar value as the final output for each image in a batch.
11
12   Class Path: AdversarialLearning -> DataModeling -> DiscriminatorDG1
13   """
14
15   def __init__(self):
16       super(DiscriminatorDG1, self).__init__()
17       self.conv_in = nn.Conv2d( 3,      64,      kernel_size=4,      stride=2,      padding=1)
18       self.conv_in2 = nn.Conv2d( 64,     128,      kernel_size=4,      stride=2,      padding=1)
19       self.conv_in3 = nn.Conv2d( 128,    256,      kernel_size=4,      stride=2,      padding=1)
20       self.conv_in4 = nn.Conv2d( 256,    512,      kernel_size=4,      stride=2,      padding=1)
21       self.conv_in5 = nn.Conv2d( 512,     1,      kernel_size=4,      stride=1,      padding=0)
22       self.bn1 = nn.BatchNorm2d(128)
23       self.bn2 = nn.BatchNorm2d(256)
24       self.bn3 = nn.BatchNorm2d(512)
25       self.sig = nn.Sigmoid()
26
27   def forward(self, x):
28       x = torch.nn.functional.leaky_relu(self.conv_in(x), negative_slope=0.2, inplace=True)
29       x = self.bn1(self.conv_in2(x))
30       x = torch.nn.functional.leaky_relu(x, negative_slope=0.2, inplace=True)
31       x = self.bn2(self.conv_in3(x))
32       x = torch.nn.functional.leaky_relu(x, negative_slope=0.2, inplace=True)
33       x = self.bn3(self.conv_in4(x))
34       x = torch.nn.functional.leaky_relu(x, negative_slope=0.2, inplace=True)
35       x = self.conv_in5(x)
36       x = self.sig(x)
37
38   return x
39
40 class GeneratorDG1(nn.Module):
41     """
42     This is an implementation of the DCGAN Generator. As was the case with the Discriminator network,
43     you again see the 4-2-1 topology here. A Generator's job is to transform a random noise
44     vector into an image that is supposed to look like it came from the training dataset. (We refer
45     to the images constructed from noise vectors in this manner as fakes.) As you will see later
46     in the "run_gan_code()" method, the starting noise vector is a 1x1 image with 100 channels. In
47     order to output 64x64 output images, the network shown below use the Transpose Convolution
48     operator nn.ConvTranspose2d with a stride of 2. If (H_in, W_in) are the height and the width
49     of the image at the input to a nn.ConvTranspose2d layer and (H_out, W_out) the same at the
50     output, the size pairs are related by
51         H_out = (H_in - 1) * s + k - 2 * p
52         W_out = (W_in - 1) * s + k - 2 * p
53
54     were s is the stride and k the size of the kernel. (I am assuming square strides, kernels, and
55     padding). Therefore, each nn.ConvTranspose2d layer shown below doubles the size of the input.
56
57     Class Path: AdversarialLearning -> DataModeling -> GeneratorDG1
58     """
59
60   def __init__(self):
61       super(GeneratorDG1, self).__init__()
62       self.latent_to_image = nn.ConvTranspose2d( 100,    512,      kernel_size=4,      stride=1,      padding=0, bias=False)
63       self.upsampler2 = nn.ConvTranspose2d( 512,    256,      kernel_size=4,      stride=2,      padding=1, bias=False)
64       self.upsampler3 = nn.ConvTranspose2d( 256,    128,      kernel_size=4,      stride=2,      padding=1, bias=False)
65       self.upsampler4 = nn.ConvTranspose2d( 128,     64,      kernel_size=4,      stride=2,      padding=1, bias=False)
66       self.upsampler5 = nn.ConvTranspose2d( 64,      3,      kernel_size=4,      stride=2,      padding=1, bias=False)
67       self.bn1 = nn.BatchNorm2d(512)
68       self.bn2 = nn.BatchNorm2d(256)
69       self.bn3 = nn.BatchNorm2d(128)
70       self.bn4 = nn.BatchNorm2d(64)
71       self.tanh = nn.Tanh()
72
73   def forward(self, x):
74       x = self.latent_to_image(x)
75       x = torch.nn.functional.relu(self.bn1(x))
76       x = self.upsampler2(x)
77       x = torch.nn.functional.relu(self.bn2(x))
78       x = self.upsampler3(x)
79       x = torch.nn.functional.relu(self.bn3(x))
80       x = self.upsampler4(x)
81       x = torch.nn.functional.relu(self.bn4(x))
82       x = self.upsampler5(x)
83       x = self.tanh(x)
84
85   return x

```

April 7, 2023

*Critic and Generator for the W-GAN. Again, mostly taken from the DLStudio implementation.*

```

83 ##### Critic-Generator CG2 #####
84 class CriticCG2(nn.Module):
85     """
86     For the sake of variety, the Critic implementation in CG2 as the same Marvin Cao's Discriminator:
87     https://github.com/caogang/wgan-gp
88
89     which in turn is the PyTorch version of the Tensorflow based Discriminator presented by the
90     authors of the paper "Improved Training of Wasserstein GANs" by Gulrajani, Ahmed, Arjovsky, Dumouli,
91     and Courville.
92
93     Class Path: AdversarialLearning -> DataModeling -> CriticCG2
94     """
95
96     def __init__(self):
97         super(CriticCG2, self).__init__()
98         self.DIM = 64
99         main = nn.Sequential(
100             nn.Conv2d(3, self.DIM, 5, stride=2, padding=2),
101             nn.ReLU(True),
102             nn.Conv2d(self.DIM, 2*self.DIM, 5, stride=2, padding=2),
103             nn.ReLU(True),
104             nn.Conv2d(2*self.DIM, 4*self.DIM, 5, stride=2, padding=2),
105             nn.ReLU(True),
106         )
107         self.main = main
108         self.output = nn.Linear(4*4*4*self.DIM, 1)
109
110     def forward(self, input):
111         input = input.view(-1, 3, 64, 64)
112         out = self.main(input)
113         out = out.view(-1, 4*4*4*self.DIM)
114         out = self.output(out)
115         out = out.mean(0)
116         out = out.view(1)
117         return out
118
119 class GeneratorCG2(nn.Module):
120     """
121     The Generator code remains the same as for DG1 shown earlier.
122
123     Class Path: AdversarialLearning -> DataModeling -> GeneratorCG2
124     """
125
126     def __init__(self):
127         super(GeneratorCG2, self).__init__()
128         self.latent_to_image = nn.ConvTranspose2d(100, 512, kernel_size=4, stride=1, padding=0, bias=False)
129         self.upsampler2 = nn.ConvTranspose2d(512, 256, kernel_size=4, stride=2, padding=1, bias=False)
130         self.upsampler3 = nn.ConvTranspose2d(256, 128, kernel_size=4, stride=2, padding=1, bias=False)
131         self.upsampler4 = nn.ConvTranspose2d(128, 64, kernel_size=4, stride=2, padding=1, bias=False)
132         self.upsampler5 = nn.ConvTranspose2d(64, 3, kernel_size=4, stride=2, padding=1, bias=False)
133         self.bn1 = nn.BatchNorm2d(512)
134         self.bn2 = nn.BatchNorm2d(256)
135         self.bn3 = nn.BatchNorm2d(128)
136         self.bn4 = nn.BatchNorm2d(64)
137         self.tanh = nn.Tanh()
138
139     def forward(self, x):
140         x = self.latent_to_image(x)
141         x = torch.nn.functional.relu(self.bn1(x))
142         x = self.upsampler2(x)
143         x = torch.nn.functional.relu(self.bn2(x))
144         x = self.upsampler3(x)
145         x = torch.nn.functional.relu(self.bn3(x))
146         x = self.upsampler4(x)
147         x = torch.nn.functional.relu(self.bn4(x))
148         x = self.upsampler5(x)
149         x = self.tanh(x)
150
151     return x
152
153 ##### Generator CG #####

```

A “mini DL studio” to input epochs, batch sizes, gpu preferences, and some other basic parameters in the experiments:

April 7, 2023

```
11
12 class MiniDLStudio(nn.Module):
13     def __init__(self, epochs, use_gpu, train_dataloader, val_dataloader, batch_size, lr, LAMBDA):
14         super(MiniDLStudio, self).__init__()
15         self.epochs = epochs
16         self.val_dataloader = val_dataloader
17         self.train_dataloader = train_dataloader
18         self.batch_size = batch_size
19         self.learning_rate = lr
20         self.LAMBDA = LAMBDA
21         if use_gpu is not None:
22             self.use_gpu = use_gpu
23             if use_gpu is True:
24                 if torch.cuda.is_available():
25                     self.device = torch.device("cuda:0")
26                 else:
27                     raise Exception("You requested GPU support, but there's no GPU on this machine")
28             else:
29                 self.device = torch.device("cpu")
30             print("assigned self.device")
31
32     def weights_init(self,m):
33         """
34             Uses the DCGAN initializations for the weights
35         """
36         classname = m.__class__.__name__
37         if classname.find('Conv') != -1:
38             nn.init.normal_(m.weight.data, 0.0, 0.02)
39         elif classname.find('BatchNorm') != -1:
40             nn.init.normal_(m.weight.data, 1.0, 0.02)
41             nn.init.constant_(m.bias.data, 0)
42
43     def calc_gradient_penalty(self, netC, real_data, fake_data):
44         """
45             Implementation by Marvin Cao: https://github.com/caogang/wgan-gp
46             Marvin Cao's code is a PyTorch version of the Tensorflow based implementation provided by
47             the authors of the paper "Improved Training of Wasserstein GANs" by Gulrajani, Ahmed,
48             Arjovsky, Dumoulin, and Courville.
49         """
50         BATCH_SIZE = self.batch_size
51         LAMBDA = self.LAMBDA
52         epsilon = torch.rand(1).cuda() #may need to change to .cuda() when using GPU
53         interpolates = epsilon * real_data + ((1 - epsilon) * fake_data)
54         interpolates = interpolates.requires_grad_(True).cuda() #may need to change to .cuda() on GPU
55         critic_interpolates = netC(interpolates)
56         gradients = torch.autograd.grad(outputs=critic_interpolates, inputs=interpolates,
57                                         grad_outputs=torch.ones(critic_interpolates.size()).cuda(), #may need to change to .cuda() later!
58                                         create_graph=True, retain_graph=True, only_inputs=True)[0]
59         gradient_penalty = ((gradients.norm(2, dim=1) - 1) ** 2).mean() * LAMBDA
60
61         return gradient_penalty
62
```

April 7, 2023

*Training and Inference Loop for the W-GAN:*

```

**
12 def run_wgan_with_gp_code(self, critic, generator, results_dir):
13     """
14         This function is meant for training a CG2-based Critic-Generator WGAN. Regarding how
15         to set the parameters of this method, see the following script in the
16         "ExamplesAdversarialLearning" directory of the distribution:
17
18             wgan_with_gp(CG2.py
19
20             """
21 #         if os.path.exists(dir_name_for_results):
22 #             files = glob.glob(dir_name_for_results + "/*")
23 #             for file in files:
24 #                 if os.path.isfile(file):
25 #                     os.remove(file)
26 #                 else:
27 #                     files = glob.glob(file + "/*")
28 #                     list(map(lambda x: os.remove(x), files))
29 #
30 #             else:
31 #                 os.mkdir(dir_name_for_results)
32 #     Set the number of channels for the 1x1 input noise vectors for the Generator:
33 nz = 100
34 netC = critic.to(self.device)
35 netG = generator.to(self.device)
36 # Initialize the parameters of the Critic and the Generator networks according to the
37 # definition of the "weights_init()" method:
38 netC.apply(self.weights_init)
39 netG.apply(self.weights_init)
40 # We will use a the same noise batch to periodically check on the progress made for the Generator:
41 fixed_noise = torch.randn(self.batch_size, nz, 1, 1, device=self.device)
42 # These are for training the Critic, 'one' is for the part of the training with actual
43 # training images, and 'minus_one' is for the part based on the images produced by the
44 # Generator:
45 one = torch.FloatTensor([1]).to(self.device)
46 minus_one = torch.FloatTensor([-1]).to(self.device)
47 # Adam optimizers for the Critic and the Generator:
48 optimizerC = optim.Adam(netC.parameters(), lr=self.learning_rate, betas=(beta1, 0.999))
49 optimizerG = optim.Adam(netG.parameters(), lr=self.learning_rate, betas=(beta1, 0.999))
50 img_list = []
51 Gen_losses = []
52 Cri_losses = []
53 iters = 0
54 gen_iterations = 0
55 start_time = time.perf_counter()
56 dataloader = self.train_dataloader
57 # For each epoch
58 for epoch in range(self.epochs):
59     data_iter = iter(dataloader)
60     i = 0
61     ncritic = 5
62     # In this version of WGAN training, we enforce the 1-Lipschitz condition on the function
63     # being learned by the Critic by requiring that the partial derivatives of the output of
64     # the Critic with respect to its input equal one in magnitude. This is referred as imposing
65     # a Gradient Penalty on the learning by the Critic. As in the previous training
66     # function, we start by turning on the "requires_grad" property of the Critic parameters:
67     while i < len(dataloader):
68         for p in netC.parameters():
69             p.requires_grad = True
70             ic = 0
71             while ic < ncritic and i < len(dataloader):
72                 ic += 1
73                 # The first two parts of what it takes to train the Critic are the same as for
74                 # a regular WGAN. We want to train the Critic to recognize the training images and,
75                 # at the same time, the Critic should try to not believe the output of the Generator.
76                 netC.zero_grad()
77                 real_images_in_batch = next(data_iter) # was data_iter.next() before
78                 i += 1
79                 real_images_in_batch = real_images_in_batch[0].to(self.device)
80                 # Need to know how many images we pulled in since at the tailend of the dataset, the
81                 # number of images may not equal the user-specified batch size:
82                 b_size = real_images_in_batch.size(0)
83                 # Note that a single scalar is produced for all the data in a batch.
84                 critic_for_reals_mean = netC(real_images_in_batch) ## this is a batch based mean
85                 # The gradient target is 'minus_one'. Note that the gradient here is one of output of
86                 # the network with respect to the learnable parameters:
87                 critic_for_reals_mean.backward(minus_one)
88                 # The second part of Critic training requires that we apply the Critic to the images

```

April 7, 2023

```

86         critic_for_reals_mean.backward(minus_one)
87         # The second part of Critic training requires that we apply the Critic to the images
88         # produced by the Generator from a fresh batch of input noise vectors.
89         noise = torch.randn(b_size, nz, 1, 1, device=self.device)
90         fakes = netG(noise)
91         # Again, a single number is produced for the whole batch:
92         critic_for_fakes_mean = netC(fakes.detach()) ## detach() returns a copy of the 'fakes' tensor that has
93         ## been removed from the computational graph. This ensures
94         ## that a subsequent call to backward() will only update the Critic
95         # The gradient target is 'one'. Note that the gradient here is one of output of
96         # the network with respect to the learnable parameters:
97         critic_for_fakes_mean.backward(one)
98         # For the third part of Critic training, we need to first estimate the Gradient Penalty
99         # of the function being learned by the Critics with respect to the input to the function.
100        gradient_penalty = self.calc_gradient_penalty(netC, real_images_in_batch, fakes)
101        gradient_penalty.backward()
102        loss_critic = critic_for_fakes_mean - critic_for_reals_mean + gradient_penalty
103        wasser_dist = critic_for_reals_mean - critic_for_fakes_mean
104        # Update the Critic
105        optimizerC.step()

106        # That brings us to the training of the Generator. First we must turn off the "requires_grad"
107        # of the Critic parameters since the Critic and the Generator are to be updated independently:
108        for p in netC.parameters():
109            p.requires_grad = False
110        netG.zero_grad()
111        # This is again a single scalar based characterization of the whole batch of the Generator images:
112        noise = torch.randn(b_size, nz, 1, 1, device=self.device)
113        fakes = netG(noise)
114        critic_for_fakes_mean = netC(fakes)
115        loss_gen = critic_for_fakes_mean
116        # The gradient target is 'minus_one'. Note that the gradient here is one of output of the network
117        # with respect to the learnable parameters:
118        loss_gen.backward(minus_one)
119        # Update the Generator
120        optimizerG.step()
121        gen_iterations += 1
122        if i % (ncritic * 20) == 0:
123            current_time = time.perf_counter()
124            elapsed_time = current_time - start_time
125            print("[epoch=%d %d %d %d] loss_critic=%7.4f loss_gen=%7.4f Wasserstein_dist=%7.4f" % (epoch+1, self
126            Gen_losses.append(loss_gen.data[0].item())
127            Cri_losses.append(loss_critic.data[0].item())
128            # Get G's output on fixed_noise for the GIF animation:
129            if (iters % 500 == 0) or ((epoch == self.epochs-1) and (i == len(self.train_dataloader)-1)):
130                with torch.no_grad():
131                    fake = netG(fixed_noise).detach().cpu() ## detach() removes the fake from comp. graph
132                    ## in order to produce a CPU compatible tensor
133                    img_list.append(torchvision.utils.make_grid(fake, padding=1, pad_value=1, normalize=True))
134            iters += 1
135
136    #RUN INFERENCE ON VAL DATA
137
138    print("about to run inference!!!")
139    with torch.no_grad():
140        count = 0
141        for n, data in enumerate(self.val_dataloader):
142            if n < 1:
143                noise = torch.randn(self.batch_size, nz, 1, 1, device=self.device)
144                fakes_inference = netG(noise)
145                print(fakes_inference.shape)
146                print("n is", n)
147
148                for num, img in enumerate(fakes_inference):
149                    if num < 5:
150                        #convert tensor to jpg, save to wgan_inference_images_fid, store in a list of "fake_paths"
151                        transform = tvt.ToPILImage()
152                        temp_jpg = transform(img)
153                        #print(type(temp_jpg))
154                        temp_num = num + (count)*self.batch_size
155                        filename = './wgan_inference_images_fid/' + str(temp_num) + '.jpg'
156                        temp_jpg.save(fp = filename)
157                        #print('saved to ', filename)
158                        fake_paths.append(filename)
159                        if num == self.batch_size-1:
160                            count = count+1
161

```

*Training and Inference Loop for BCE-GAN:*

April 7, 2023

```

3 fake_paths_dcgan = []
4 def run_gan_code(self, discriminator, generator, results_dir):
5     """
6         This function is meant for training a Discriminator-Generator based Adversarial Network.
7         The implementation shown uses several programming constructs from the "official" DCGAN
8         implementations at the PyTorch website and at GitHub.
9
10    Regarding how to set the parameters of this method, see the following script
11
12        dcgan_DG1.py
13
14    in the "ExamplesAdversarialLearning" directory of the distribution.
15    """
16
17    dir_name_for_results = results_dir
18    # if os.path.exists(dir_name_for_results):
19    #     files = glob.glob(dir_name_for_results + "/*")
20    #     for file in files:
21    #         if os.path.isfile(file):
22    #             os.remove(file)
23    #     else:
24    #         files = glob.glob(file + "/*")
25    #         list(map(lambda x: os.remove(x), files))
26    # else:
27    #     os.mkdir(dir_name_for_results)
28
29    # Set the number of channels for the 1x1 input noise vectors for the Generator:
30    nz = 100
31    netD = discriminator.to(self.device)
32    netG = generator.to(self.device)
33
34    # Initialize the parameters of the Discriminator and the Generator networks according to the
35    # definition of the "weights_init()" method:
36    netD.apply(self.weights_init)
37    netG.apply(self.weights_init)
38
39    # We will use a the same noise batch to periodically check on the progress made for the Generator:
40    fixed_noise = torch.randn(self.batch_size, nz, 1, 1, device=self.device)
41
42    # Establish convention for real and fake labels during training
43    real_label = 1
44    fake_label = 0
45
46    # Adam optimizers for the Discriminator and the Generator:
47    optimizerD = optim.Adam(netD.parameters(), lr=self.learning_rate, betas=(adversarial_beta, 0.999))
48    optimizerG = optim.Adam(netG.parameters(), lr=self.learning_rate, betas=(adversarial_beta, 0.999))
49
50    # Establish the criterion for measuring the loss at the output of the Discriminator network:
51    criterion = nn.BCELoss()
52
53    # We will use these lists to store the results accumulated during training:
54    img_list = []
55    G_losses = []
56    D_losses = []
57    iters = 0
58
59    print("\n\nStarting Training Loop...\n\n")
60    start_time = time.perf_counter()
61    for epoch in range(self.epochs):
62        g_losses_per_print_cycle = []
63        d_losses_per_print_cycle = []
64
65        # For each batch in the dataloader
66        for i, data in enumerate(self.train_dataloader, 0):
67            #print("starting new batch")
68            ## Maximization Part of the Min-Max Objective of Eq. (3):
69            ##
70            ## As indicated by Eq. (3) in the DCGAN part of the doc section at the beginning of this
71            ## file, the GAN training boils down to carrying out a min-max optimization. Each iterative
72            ## step of the max part results in updating the Discriminator parameters and each iterative
73            ## step of the min part results in the updating of the Generator parameters. For each
74            ## batch of the training data, we first do max and then do min. Since the max operation
75            ## affects both terms of the criterion shown in the doc section, it has two parts: In the
76            ## first part we apply the Discriminator to the training images using 1.0 as the target;
77            ## and, in the second part, we supply to the Discriminator the output of the Generator
78            ## and use 0 as the target. In what follows, the Discriminator is being applied to
79            ## the training images:
80            netD.zero_grad()
81            #real_images_in_batch = data[0].to(self.device)
82            real_images_in_batch = data.to(self.device)
83            #print("real_images_in_batch is:", real_images_in_batch.shape)
84            #print("data shape is", data.shape)
85            # Need to know how many images we pulled in since at the tailend of the dataset, the
86            # number of images may not equal the user-specified batch size:
87            b_size = real_images_in_batch.size(0)
88            #print("b_size is", b_size)
89            label = torch.full((b_size,), real_label, dtype=torch.float, device=self.device)

```

April 7, 2023

```

85      ## That brings us the second part of what it takes to carry out the max operation on the
86      ## min-max criterion shown in Eq. (3) in the doc section at the beginning of this file.
87      ## part calls for applying the Discriminator to the images produced by the Generator from
88      ## noise:
89      noise = torch.randn(b_size, nz, 1, 1, device=self.device)
90      fakes = netG(noise)
91      label.fill_(fake_label)
92      ## The call to fakes.detach() in the next statement returns a copy of the 'fakes' tensor
93      ## that does not exist in the computational graph. That is, the call shown below first
94      ## makes a copy of the 'fakes' tensor and then removes it from the computational graph.
95      ## The original 'fakes' tensor continues to remain in the computational graph. This ploy
96      ## ensures that a subsequent call to backward() in the 3rd statement below would only
97      ## update the netD weights.
98      output = netD(fakes.detach()).view(-1)
99      lossD_for_fakes = criterion(output, label)
100     ## At this point, we do not care if the following call also calculates the gradients
101     ## wrt the Discriminator weights since we are going to next iteration with 'netD.zero_grad()':
102     lossD_for_fakes.backward()
103     lossD = lossD_for_reals + lossD_for_fakes
104     d_losses_per_print_cycle.append(lossD)
105     ## Only the Discriminator weights are incremented:
106     optimizerD.step()
107
108    ## Minimization Part of the Min-Max Objective of Eq. (3):
109    ##
110    ## That brings to the min part of the max-min optimization described in Eq. (3) the doc
111    ## section at the beginning of this file. The min part requires that we minimize
112    ## "1 - D(G(z))" which, since D is constrained to lie in the interval (0,1), requires that
113    ## we maximize D(G(z)). We accomplish that by applying the Discriminator to the output
114    ## of the Generator and use 1 as the target for each image:
115    netG.zero_grad()
116    label.fill_(real_label)
117    output = netD(fakes).view(-1)
118    lossG = criterion(output, label)
119    g_losses_per_print_cycle.append(lossG)
120    lossG.backward()
121    ## Only the Generator parameters are incremented:
122    optimizerG.step()
123
124    if i % 100 == 99:
125        current_time = time.perf_counter()
126        elapsed_time = current_time - start_time
127        mean_D_loss = torch.mean(torch.FloatTensor(d_losses_per_print_cycle))
128        mean_G_loss = torch.mean(torch.FloatTensor(g_losses_per_print_cycle))
129        print("[epoch=%d/%d iter=%4d elapsed_time=%5d secs] mean_D_loss=%7.4f mean_G_loss=%7.4f" %
130              (epoch+1, self.epochs, (i+1), elapsed_time, mean_D_loss, mean_G_loss))
131        d_losses_per_print_cycle = []
132        g_losses_per_print_cycle = []
133        G_losses.append(lossG.item())
134        D_losses.append(lossD.item())
135        if (iters % 500 == 0) or ((epoch == self.epochs-1) and (i == len(self.train_dataloader)-1)):
136            with torch.no_grad():
137                fake = netG(fixed_noise).detach().cpu() ## detach() removes the fake from comp. graph.
138                                            ## for creating its CPU compatible version
139                img_list.append(torchvision.utils.make_grid(fake, padding=1, pad_value=1, normalize=True))
140        iters += 1
...

```

April 7, 2023

```

142     #run inference using trained generator:
143     with torch.no_grad():
144         count = 0
145         for n, data in enumerate(val_dataloader):
146             #if n < 1:
147                 data_shp = data.shape[0]
148                 print(data_shp)
149                 noise = torch.randn(data_shp, nz, 1, 1, device=self.device)
150                 fakes_inference = netG(noise)
151                 print("fakes_inference shape", fakes_inference.shape)
152                 #print("n is", n)
153
154             for num, img in enumerate(fakes_inference):
155                 #if num < 5:
156                     #convert tensor to jpg, save to wgan_inference_images_fid, store in a list of "fake_paths"
157                     transform = tvt.ToPILImage()
158                     temp_jpg = transform(img)
159                     #print(type(temp_jpg))
160                     temp_num = num + (count)*self.batch_size
161                     filename = './dcgan_inference_images_fid/' + str(temp_num) + '.jpg'
162                     temp_jpg.save(fp = filename)
163                     #print('saved to ', filename)
164                     fake_paths_dcgan.append(filename)
165                     if num == self.batch_size-1:
166                         count = count+1
167
168             print("fake paths is", fake_paths_dcgan, "and has \n", len(fake_paths_dcgan), "images")
169
170
171

```

*Code to kick off training and inference for both networks:*

```

1 ##### CODE TO KICK OFF DCGAN TRAINING/EVAL #####
2
3 mdls = MiniDLStudio(epochs = 100, use_gpu = True, train_dataloader = train_dataloader, val_dataloader = val_dataloader, batch_size=64,
4                      lr = 1e-5, LAMBDA = LAMBDA)
5 run_gan_code(mdls, discriminator, generator, results_dir)

```

```

1
2 #run wgan code:
3 critic = CriticCG2()
4 generator = GeneratorCG2()
5 print(os.getcwd())
6 learning_rate = 0.00001
7 #wgan.show_sample_images_from_dataset(dls)
8 mdls_wgan = MiniDLStudio(epochs = 100, use_gpu = True, train_dataloader = train_dataloader, val_dataloader = val_dataloader, batch_size=40,
9                          lr = learning_rate, LAMBDA = LAMBDA)
10 run_wgan_with_gp_code(mdls_wgan, critic=critic, generator=generator,
11                        results_dir=results_dir_wgan_gp)
12
13

```

*Code to Compute FID scores (note that for both networks they are very similar, so, only one is pasted):*

April 7, 2023

```

1 ### compute FID for DC GAN ###
2 #first, test opening fake paths:
3 PIL_img = Image.open('./dcgan_inference_images_fid/0.jpg')
4 display(PIL_img)
5
6 #put real validation images in list:
7 real_paths_dcgan = []
8 d = os.listdir('./pizzas/eval')
9 print(len(d))
10 for n, i in enumerate(d):
11     real_paths_dcgan.append('./pizzas/eval/' + str(i))
12
13 from pytorch_fid.fid_score import calculate_activation_statistics
14 from pytorch_fid.fid_score import calculate_frechet_distance
15 from pytorch_fid.inception import InceptionV3
16
17 dims = 2048
18 block_idx = InceptionV3.BLOCK_INDEX_BY_DIM[dims]
19 model = InceptionV3([block_idx]).to('cuda:0')
20 m1, s1 = calculate_activation_statistics(real_paths_dcgan, model, device='cuda:0')
21 m2, s2 = calculate_activation_statistics(fake_paths_dcgan, model, device='cuda:0')
22 fid_value = calculate_frechet_distance(m1, s1, m2, s2)
23 print(fid_value)

```

*Code to plot generated images:*

```

1 #generate 4 x 4 plot of generated images:
2 fig, ax = plt.subplots(4, 4, figsize = (8,8))
3 for i in range(4):
4     for j in range(4):
5         temp_num = np.random.randint(low=0, high=999, size=None, dtype=int)
6         PIL_img = Image.open('./dcgan_inference_images_fid/' + str(temp_num) + '.jpg')
7         ax[i,j].imshow(np.uint8(PIL_img))
8

```