

In this lab, we aim to predict whether product reviews from Amazon have a positive or negative sentiment about the product. The neural network architecture is a recurrent neural network, specifically, a Gated Recurrent Network (GRU) that keeps feeding into itself outputs from the previous run. The lab gives details on the exact prediction mechanism a GRU uses. In this lab, three models are written, trained, and tested. First, a GRU is built from scratch and then wrapped in a recurrent network. Second, PyTorch's default GRU is used to predict sentiments. Lastly, PyTorch's bi-directional GRU is used, which adds a boost to the prediction score on the same dataset.

Part 1: Sentiment Analysis with Own GRU

Product reviews were put into PyTorch's dataloader utility using the DLStudio code. The review words are “translated” to 300 dimensional embedding vectors, and each review is loaded as a tensor with the first axis representing the number of words in the review and the second axis representing the embedding dimension (300). All the following results will use this same format for loading data.

Code for own GRU, logic, and mitigating vanishing gradients:

This part of the lab required implementing our own GRU logic instead of using the module provided by PyTorch. The gated unit logic, specifically, the computation of the hidden gate (z_t), reset gate (r_t), the candidate hidden state ($h_{t\sim}$), and the new hidden state (h_t) is below:

```

3  class GruNet(nn.Module):
4      def __init__(self, input_size, hidden_size):
5          super(GruNet, self).__init__()
6          self.input_size = input_size
7          self.hidden_size = hidden_size
8
9          #weights to create the hidden gate, zt
10         self.fc_hidden_gate = nn.Linear(input_size + hidden_size, hidden_size)
11
12         #weights to create the reset gate
13         self.fc_reset_gate = nn.Linear(input_size + hidden_size, hidden_size)
14
15         #weights to create the candidate hidden state
16         self.fc_hidden_state = nn.Linear(input_size + hidden_size, hidden_size)
17
18         self.sigmoid = nn.Sigmoid()
19         self.tanh = nn.Tanh()
20
21
22
23     def forward(self, x, ht_minus1):
24         #create zt, the hidden gate
25         concat_xt_ht_minus1 = torch.cat((x, ht_minus1), dim = 0)
26         pre_act_update_gate = self.fc_hidden_gate(concat_xt_ht_minus1)
27         zt = sigmoid(pre_act_update_gate) #zt, update gate, created
28
29         #create rt, the reset gate
30         pre_act_reset_gate = self.fc_reset_gate(concat_xt_ht_minus1)
31         rt = sigmoid(pre_act_reset_gate)
32         #print(rt.shape)
33
34         #create ht_tilde, the candidate hidden state
35         pre_act_hidden_state = self.fc_hidden_state(concat_xt_ht_minus1)
36         ht_candidate = self.tanh(pre_act_hidden_state)
37         #print(ht_candidate.shape)
38
39         #create new hidden state, which weights 1 - hidden gate element wise multiplied by prev hidden state
40         # added to hidden gate element wise multiplied by the candidate hidden
41         ht_new = (((1 - zt) * ht_minus1) + (zt * ht_candidate))
42
43     return ht_new

```

The hidden gate is computed by first concatenating x_t and h_{t-1} , resulting in a 1×400 vector. This 1×400 vector is multiplied by a weight matrix that is 400×100 , resulting in a 1×100 vector. After a sigmoid is applied to that vector, we have the resulting z_t . In a very similar manner, r_t , the reset gate is computed, resulting in another 1×100 vector. This reset gate is factored into computing a candidate hidden state ($h_{t'}$). Now, both the hidden gate (z_t) and the candidate hidden state ($h_{t'}$) are accessible in the computation of the actual new hidden state in the form of $(1 - z_t) * h_{t-1} + z_t * h_{t'}$. We can see that if z_t is close to one, then, the old hidden state (h_{t-1}) will be replaced by the candidate hidden state ($h_{t'}$). If z_t is close to 0, then the old hidden state (h_{t-1}) will dominate the output, and it would be like the new word has not contributed much to learning of the semantic meaning (sentiment, in this case) of the review. This mechanism basically obviates the need to hold many hidden states and create an extremely high layer network, which is how the problem of vanishing gradients is somewhat mitigated in the GRU implementation.

Then, the recurrent network is basically akin to using the new hidden state as input to the GRU network posted above. This is handled in another RNN class, where there is a loop in the forward model to produce and output a two dimensional vector after the last hidden state and word embedding has been passed to the GRU. This recurrent network is posted below:

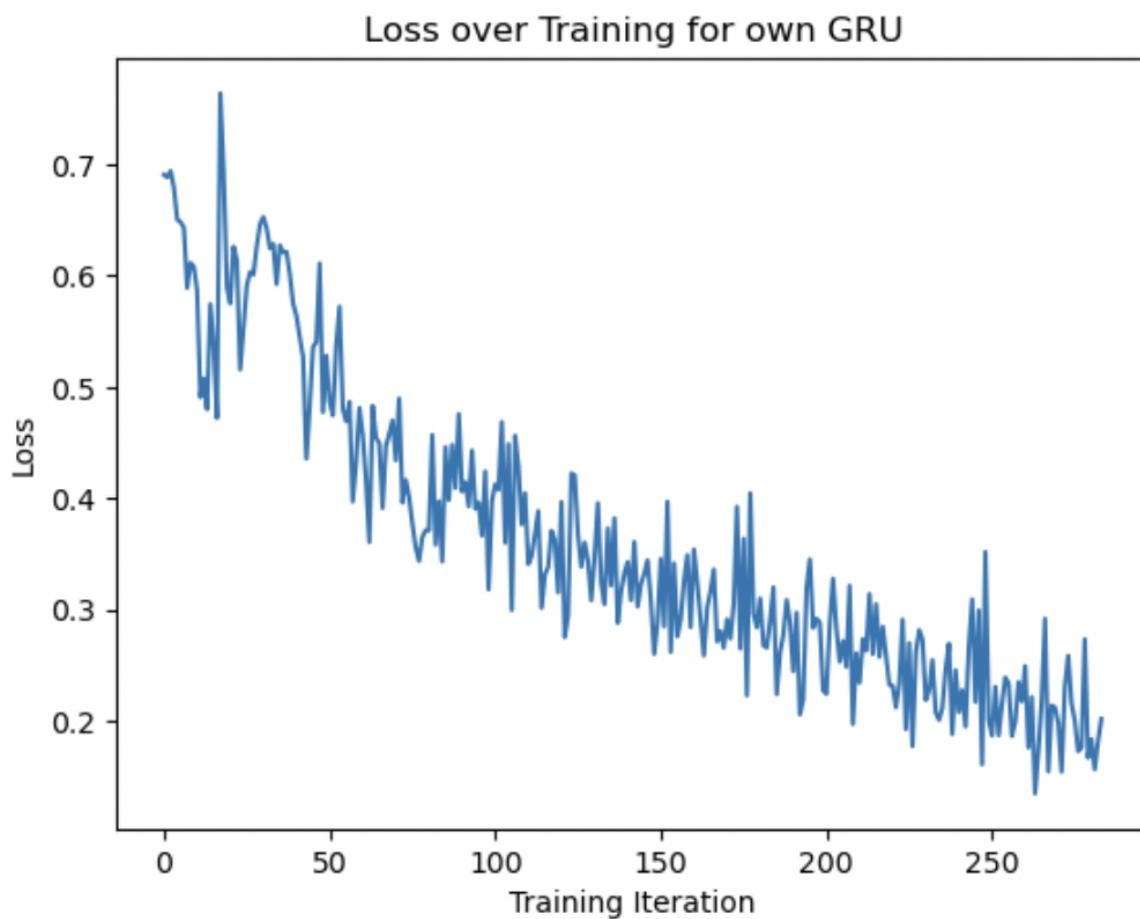
```

45 class RNN(nn.Module):
46
47     def __init__(self, input_size, hidden_size):
48         super(RNN, self).__init__()
49         self.gru = GruNet(input_size, hidden_size) #instantiate GRU
50         self.fc = nn.Linear(100, 2)
51         self.relu = nn.ReLU()
52         self.logsoftmax = nn.LogSoftmax(dim = 0)
53
54     def forward(self, x):
55         for i, word_emb in enumerate(range(x.shape[0])): #this loop runs GRU for each embedding and hidden vector
56             if i == 0: #no hidden state yet, init to 0:
57                 ht_new = torch.zeros(100)
58                 ht_new = self.gru(x[word_emb], ht_new)
59
60             else:
61                 ht_new = self.gru(x[word_emb], ht_new)
62
63         out = self.fc(self.relu(ht_new)) #take last hidden state vector to output size (2)
64         out = self.logsoftmax(out) #apply logsoftmax activation (before computing NLL Loss outside of the class)
65
66     return out
67

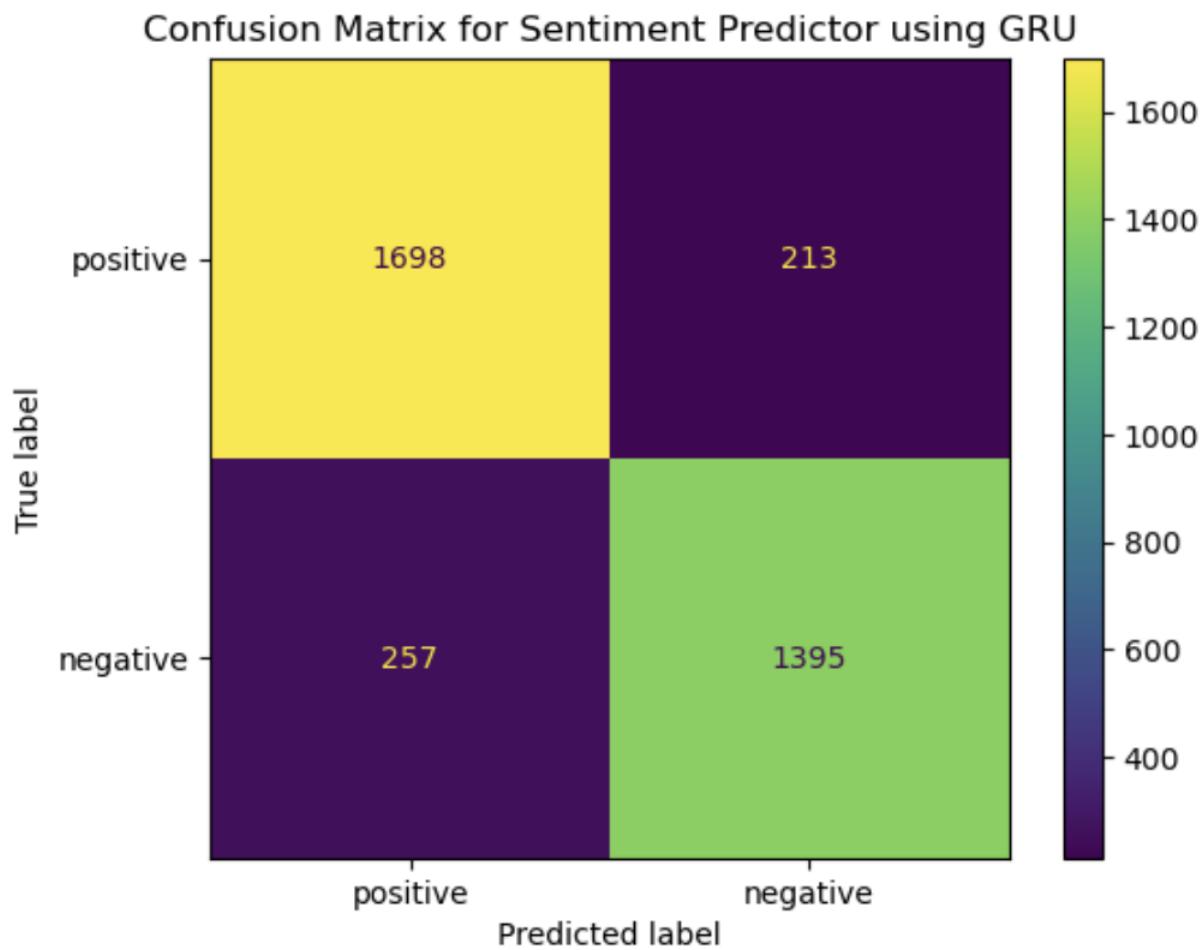
```

Results on own GRU:

Training loss progression:



Confusion Matrix for own GRU:



Total Accuracy for own GRU:

total acc for pos:

0.8885400313971743

total accuracy: 0.868088689306764

total acc for neg:

0.8444309927360775

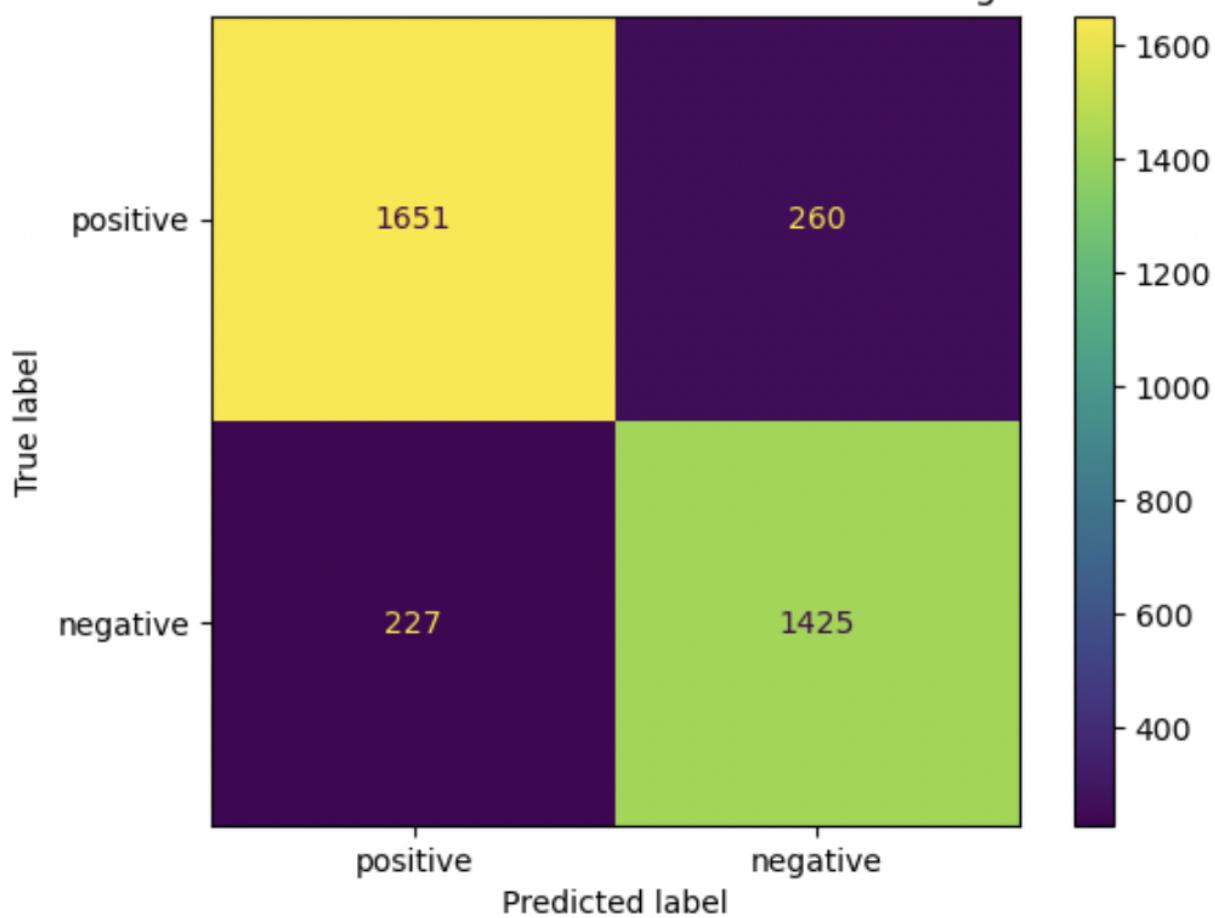
Part 2: Sentiment Analysis using torch.nn.GRU

Training Loss curve for torch.nn.GRU (bidirectional=False)



Confusion Matrix for nn.GRU (non-bidirectional variant)

Confusion Matrix for Sentiment Predictor using GRU



Overall Accuracy for nn.GRU:

total acc for pos:

0.8639455782312925

total accuracy: 0.8633174291327533

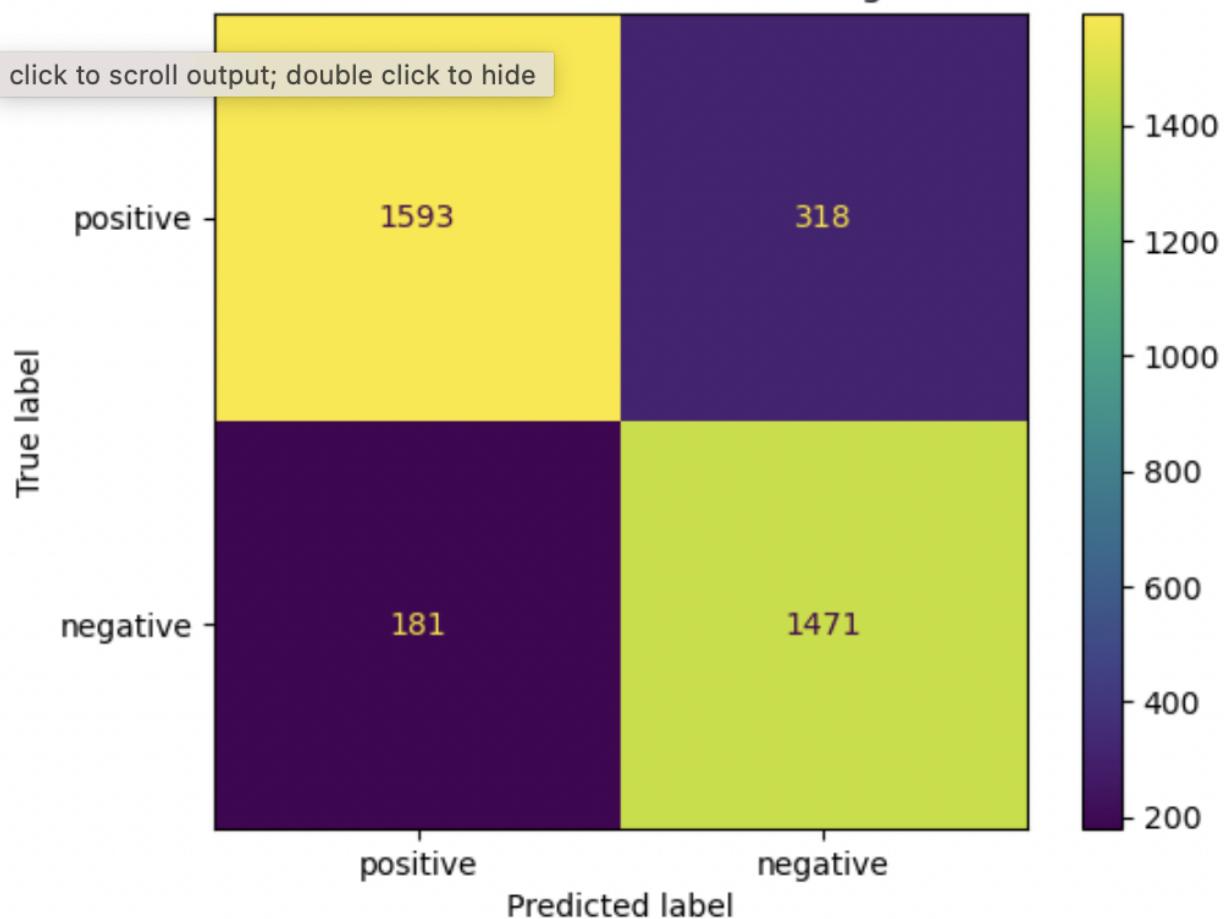
total acc for neg:

0.862590799031477

Part 3: Sentiment Analysis using nn.GRU bidirectional



Confusion Matrix for Sentiment Predictor using Bidirectional GRU



Overall Accuracy for nn.GRU bidirectional:

total acc for pos:

0.8335949764521193

total accuracy: 0.8599494807746281

total acc for neg:

0.8904358353510896

Part 4: Comparison of all RNNs

Total Accuracy Comparison		
Own GRU	nn.GRU	nn.GRU, bidirectional
86.8 %	86.3%	85.9%

Every network was trained for four epochs. Overall, all the RNNs achieve an extremely similar total accuracy of around 86%. If one cares most about overall accuracy in identifying positive sentiments correctly, the bidirectional variant would be the best while if one cares most about correctly identifying negative sentiments, the hand crafted RNN may be best. However, the results are so similar that multiple test runs should be run to truly understand if there is a model that sticks out as the “best” RNN.

Part 5: Listing of Source Code

Extracting embeddings via dataloader. Note, most of this is done with the code provided in DL Studio

April 16, 2023

```
1 #Implement logic to get the Word2Vec embeddings:
2 # Use DL studio to get sentiment data we want to train:
3 import gzip
4 import os, sys
5 import pickle
6 print(os.getcwd())
7 dataset_archive_train = "data/sentiment_dataset_train_400.tar.gz"
8 dataset_archive_test = "data/sentiment_dataset_test_400.tar.gz"
9 path_to_saved_embeddings = "word2vec_saved_embeddings/"
10 dataroot = "/Users/alim/Documents/ECE60146/hw8/"
11
12
13
14
15 class MiniDLStudio():
16     def __init__(self, dataroot, path_saved_model, momentum, learning_rate,
17                  epochs, batch_size, classes, use_gpu):
18         self.dataroot = dataroot
19         self.path_saved_model = path_saved_model
20         self.momentum = momentum
21         self.learning_rate = learning_rate
22         self.epochs = epochs
23         self.batch_size = batch_size
24         self.classes = classes
25         self.use_gpu = use_gpu
26
27 dls = MiniDLStudio(
28                     dataroot = dataroot,
29                     path_saved_model = "./saved_model",
30                     momentum = 0.9,
31                     learning_rate = 1e-5,
32                     epochs = 1,
33                     batch_size = 1,
34                     classes = ('negative','positive'),
35                     use_gpu = False
36                 )
37
38
```

April 16, 2023

```
1 class TextClassificationWithEmbeddings(nn.Module):
2     """
3         The text processing class described previously, TextClassification, was based on
4         using one-hot vectors for representing the words. The main challenge we faced
5         with one-hot vectors was that the larger the size of the training dataset, the
6         larger the size of the vocabulary, and, therefore, the larger the size of the
7         one-hot vectors. The increase in the size of the one-hot vectors led to a
8         model with a significantly larger number of learnable parameters --- and, that,
9         in turn, created a need for a still larger training dataset. Sounds like a classic
10        example of a vicious circle. In this section, I use the idea of word embeddings
11        to break out of this vicious circle.
12
13    Word embeddings are fixed-sized numerical representations for words that are
14    learned on the basis of the similarity of word contexts. The original and still
15    the most famous of these representations are known as the word2vec
16    embeddings. The embeddings that I use in this section consist of pre-trained
17    300-element word vectors for 3 million words and phrases as learned from Google
18    News reports. I access these embeddings through the popular Gensim library.
19
20    Class Path: DLStudio -> TextClassificationWithEmbeddings
21    """
22    def __init__(self, dl_studio, dataserver_train=None, dataserver_test=None, dataset_file_train=None, dataset_f
23        super(TextClassificationWithEmbeddings, self).__init__()
24        self.dl_studio = dl_studio
25        self.dataserver_train = dataserver_train
26        self.dataserver_test = dataserver_test
27
28    class SentimentAnalysisDataset(torch.utils.data.Dataset):
29        """
30            In relation to the SentimentAnalysisDataset defined for the TextClassification section of
31            DLStudio, the __getitem__() method of the dataloader must now fetch the embeddings from
32            the word2vec word vectors.
33
34            Class Path: DLStudio -> TextClassificationWithEmbeddings -> SentimentAnalysisDataset
35        """
36        def __init__(self, dl_studio, train_or_test, dataset_file, path_to_saved_embeddings=None):
37            super(TextClassificationWithEmbeddings.SentimentAnalysisDataset, self).__init__()
38            import gensim.downloader as gen_api
39            # self.word_vectors = gen_api.load("word2vec-google-news-300")
40            self.path_to_saved_embeddings = path_to_saved_embeddings
41            self.train_or_test = train_or_test
42            root_dir = dl_studio.dataroot
43            f = gzip.open(root_dir + dataset_file, 'rb')
```

April 16, 2023

```
44
45     dataset = f.read()
46     if path_to_saved_embeddings is not None:
47         import gensim.downloader as genapi
48         from gensim.models import KeyedVectors
49         if os.path.exists(path_to_saved_embeddings + 'vectors.kv'):
50             self.word_vectors = KeyedVectors.load(path_to_saved_embeddings + 'vectors.kv')
51         else:
52             print("""\n\nSince this is your first time to install the word2vec embeddings, it may take
53             """'\na couple of minutes. The embeddings occupy around 3.6GB of your disk space.\n\n")
54             self.word_vectors = genapi.load("word2vec-google-news-300")
55             ## 'kv' stands for "KeyedVectors", a special datatype used by gensim because it
56             ## has a smaller footprint than dict
57             self.word_vectors.save(path_to_saved_embeddings + 'vectors.kv')
58     if train_or_test == 'train':
59         if sys.version_info[0] == 3:
60             self.positive_reviews_train, self.negative_reviews_train, self.vocab = pickle.loads(dataset)
61         else:
62             self.positive_reviews_train, self.negative_reviews_train, self.vocab = pickle.loads(dataset)
63             self.categories = sorted(list(self.positive_reviews_train.keys()))
64             self.category_sizes_train_pos = {category : len(self.positive_reviews_train[category]) for category in self.categories}
65             self.category_sizes_train_neg = {category : len(self.negative_reviews_train[category]) for category in self.categories}
66             self.indexed_dataset_train = []
67             for category in self.positive_reviews_train:
68                 for review in self.positive_reviews_train[category]:
69                     self.indexed_dataset_train.append([review, category, 1])
70             for category in self.negative_reviews_train:
71                 for review in self.negative_reviews_train[category]:
72                     self.indexed_dataset_train.append([review, category, 0])
73             random.shuffle(self.indexed_dataset_train)
74     elif train_or_test == 'test':
75         if sys.version_info[0] == 3:
76             self.positive_reviews_test, self.negative_reviews_test, self.vocab = pickle.loads(dataset),
77         else:
78             self.positive_reviews_test, self.negative_reviews_test, self.vocab = pickle.loads(dataset)
79             self.vocab = sorted(self.vocab)
80             self.categories = sorted(list(self.positive_reviews_test.keys()))
81             self.category_sizes_test_pos = {category : len(self.positive_reviews_test[category]) for category in self.categories}
82             self.category_sizes_test_neg = {category : len(self.negative_reviews_test[category]) for category in self.categories}
83             self.indexed_dataset_test = []
84             for category in self.positive_reviews_test:
85                 for review in self.positive_reviews_test[category]:
86                     self.indexed_dataset_test.append([review, category, 1])
87             for category in self.negative_reviews_test:
88                 for review in self.negative_reviews_test[category]:
89                     self.indexed_dataset_test.append([review, category, 0])
```

```

97         else:
98             next
99             review_tensor = torch.FloatTensor( list_of_embeddings )
100            return review_tensor
101
102    def sentiment_to_tensor(self, sentiment):
103        """
104            Sentiment is ordinarily just a binary valued thing. It is 0 for negative
105            sentiment and 1 for positive sentiment. We need to pack this value in a
106            two-element tensor.
107        """
108        sentiment_tensor = torch.zeros(2)
109        if sentiment == 1:
110            sentiment_tensor[1] = 1
111        elif sentiment == 0:
112            sentiment_tensor[0] = 1
113        sentiment_tensor = sentiment_tensor.type(torch.long)
114        return sentiment_tensor
115
116    def __len__(self):
117        if self.train_or_test == 'train':
118            return len(self.indexed_dataset_train)
119        elif self.train_or_test == 'test':
120            return len(self.indexed_dataset_test)
121
122    def __getitem__(self, idx):
123        sample = self.indexed_dataset_train[idx] if self.train_or_test == 'train' else self.indexed_dataset_test
124        review = sample[0]
125        review_category = sample[1]
126        review_sentiment = sample[2]
127        review_sentiment = self.sentiment_to_tensor(review_sentiment)
128        review_tensor = self.review_to_tensor(review)
129        category_index = self.categories.index(review_category)
130        sample = {'review' : review_tensor,
131                  'category' : category_index, # should be converted to tensor, but not yet used
132                  'sentiment' : review_sentiment }
133        return sample
134
135    def load_SentimentAnalysisDataset(self, dataserver_train, dataserver_test ):
136        self.train_dataloader = torch.utils.data.DataLoader(dataserver_train,
137                                              batch_size=self.dl_studio.batch_size,shuffle=True, num_workers=2) #wraps the dataset in this
138        self.test_dataloader = torch.utils.data.DataLoader(dataserver_test,
139                                              batch_size=self.dl_studio.batch_size,shuffle=False, num_workers=2)
140
141
142    def load_SentimentAnalysisDataset(self, dataserver_train, dataserver_test ):
143        self.train_dataloader = torch.utils.data.DataLoader(dataserver_train,
144                                              batch_size=self.dl_studio.batch_size,shuffle=True, num_workers=2) #wraps the dataset in this
145        self.test_dataloader = torch.utils.data.DataLoader(dataserver_test,
146                                              batch_size=self.dl_studio.batch_size,shuffle=False, num_workers=2)
147
148    text_cl = TextClassificationWithEmbeddings(dl_studio = dls)
149    dataserver_train = text_cl.SentimentAnalysisDataset(dl_studio = dls, train_or_test = 'train',
150                                                       dataset_file = dataset_archive_train,
151                                                       path_to_saved_embeddings = path_to_saved_embeddings)
152
153    dataserver_test = TextClassificationWithEmbeddings.SentimentAnalysisDataset(
154        train_or_test = 'test',
155        dl_studio = dls,
156        dataset_file = dataset_archive_test,
157        path_to_saved_embeddings = path_to_saved_embeddings
158    )
159    text_cl.dataserver_train = dataserver_train #assigns the dataserver_train to self.dataserver_train
160    text_cl.dataserver_test = dataserver_test #assigns the dataserver_test to self.dataserver_test
161
162    text_cl.load_SentimentAnalysisDataset(dataserver_train, dataserver_test)
163
164

```

Training and Eval Code for Own GRU:

```

1 #Training Loop for handcrafted GRU:
2 #GRUNet = GruNet(input_size = 300, hidden_size = 100)
3 rnn = RNN(input_size = 300, hidden_size = 100)
4 epochs = 4
5 criterion = nn.NLLLoss()
6 optimizer = optim.Adam(rnn.parameters(), lr=dls.learning_rate)
7 running_loss = 0.0
8 training_loss_tally_ownGRU = []
9 start_time = time.perf_counter()
10 for epoch in range(epochs):
11     for n, data in enumerate(text_cl.dataserver_train):
12         review = data['review'] #review is a tensor of embeddings
13         sentiment = data['sentiment'] #sentiment is a tensor of GT sentiment:
14         out = rnn(review)
15         target = torch.argmax(sentiment)
16         loss = criterion(out, target)
17         running_loss += loss.item()
18         loss.backward()
19         optimizer.step()
20         if n % 200 == 199:
21             avg_loss = running_loss / float(200)
22             training_loss_tally_ownGRU.append(avg_loss)
23             current_time = time.perf_counter()
24             time_elapsed = current_time - start_time
25             print("[epoch:%d iter:%d elapsed_time:%d secs]      loss: %.5f" % (epoch+1,n+1, time_elapsed,avg_loss))
26             running_loss = 0.0
27
28
1 ## Prediction / Eval Code for my own GRU:
2 correct_pos_sentiment = 0.0
3 total_pos_sentiment = 0.0
4 correct_neg_sentiment = 0.0
5 total_neg_sentiment = 0.0
6 sentiment_pred = []
7 sentiment_label = []
8 mapping = { 0 : 'negative', 1: 'positive'}
9 with torch.no_grad():
10     for n, data in enumerate(text_cl.dataserver_test):
11
12         #print('n is', n)
13         sentiment = data['sentiment']
14         review = data['review']
15         out = rnn(review) #goes through one sequence of embeddings (i.e one review)
16
17         GT = torch.argmax(sentiment).item()
18         sentiment_label.append(mapping[GT])
19         pred = torch.argmax(out).item()
20         sentiment_pred.append(mapping[pred])
21         if (GT == 1): #compute positive sentiment accuracy
22             total_pos_sentiment += 1
23             if pred == GT:
24                 #print('accurate pos sentiment pred')
25                 correct_pos_sentiment += 1
26
27         if (GT == 0): #compute negative sentiment accuracy
28             total_neg_sentiment += 1
29             if pred == GT:
30                 #print('accurate neg sentiment pred')
31                 correct_neg_sentiment += 1
32
33
34     print("total acc for pos:      total acc for neg:")
35     print(correct_pos_sentiment / total_pos_sentiment, ' ', correct_neg_sentiment / total_neg_sentiment)
36     print('total accuracy:', (correct_pos_sentiment + correct_neg_sentiment) / len(sentiment_label))
37
38 from sklearn.metrics import confusion_matrix
39 y_true = sentiment_label
40 y_pred = sentiment_pred
41 print(len(y_true), len(y_pred))
42 confusion_matrix=confusion_matrix(y_true, y_pred, labels = [ "positive", "negative"])
43 disp = ConfusionMatrixDisplay(confusion_matrix, display_labels = [ "positive", "negative"])
44 disp.plot()
45 disp.ax_.set_title("Confusion Matrix for Sentiment Predictor using own GRU")
46 plt.show()
47

```

Network, Training and Eval Code for nn.GRU:

```
1 class GRUNetPyTorch(nn.Module):
2     """
3         For this embeddings adapted version of the GRU net shown earlier, we can assume that
4         the 'input_size' for a tensor representing a word is always 300.
5         Source: https://blog.floydhub.com/gru-with-pytorch/
6         with the only modification that the final output of forward() is now
7         routed through LogSoftmax activation.
8
9     Class Path: DLStudio -> TextClassificationWithEmbeddings -> GRUNetWithEmbeddings
10    """
11    def __init__(self, input_size, hidden_size, output_size, num_layers=1):
12        """
13            -- input_size is the size of the tensor for each word in a sequence of words. If you word2vec
14                embedding, the value of this variable will always be equal to 300.
15            -- hidden_size is the size of the hidden state in the RNN
16            -- output_size is the size of output of the RNN. For binary classification of
17                input text, output_size is 2.
18            -- num_layers creates a stack of GRUs
19        """
20        super(GRUNetPyTorch, self).__init__()
21        self.input_size = input_size
22        self.hidden_size = hidden_size
23        self.num_layers = num_layers
24        self.output_size = output_size
25        self.gru = nn.GRU(self.input_size, self.hidden_size, self.num_layers)
26        self.fc = nn.Linear(self.hidden_size, self.output_size)
27        self.relu = nn.ReLU()
28        self.logsoftmax = nn.LogSoftmax(dim=1)
29
30    def forward(self, x, h):
31        #print("in input: x shp and then h shape:", x.shape, h.shape)
32        out, h = self.gru(x, h)
33        #print("after running through self.gru:", out.shape, h.shape)
34        #print("out[:, -1].shape", out[:, -1].shape)
35        out = self.fc(self.relu(out[:, -1]))
36
37        #print("after relu applied to out[:, -1] and a fc(hidden size --> output size):", out.shape)
38        out = self.logsoftmax(out)
39        #print("output after logsoftmax:", out.shape)
40        return out, h
41
42    def init_hidden(self):
43        weight = next(self.parameters()).data
44        #           num_layers   batch_size   hidden_size
45        hidden = weight.new( 2,           1,           self.hidden_size ).zero_()
46        return hidden
47
```

April 16, 2023

```
1 model = GRUNetPyTorch(input_size = 300, hidden_size = 100, output_size = 2, num_layers=2)
2 epochs = 4
3 criterion = nn.NLLLoss()
4 optimizer = optim.Adam(model.parameters(), lr=dls.learning_rate)
5 running_loss = 0.0
6 training_loss_tally_nnGRU = []
7 start_time = time.perf_counter()
8
9 for epoch in range(epochs):
10     for n, data in enumerate(text_cl.dataserver_train):
11         review = data['review']
12         sentiment = data['sentiment']
13         h = model.init_hidden()
14         #print('h.shape:',h.shape)
15         #input to an instance of torch.nn.GRU = tensor with shape seq len, batchsize, input_size
16
17         review = torch.unsqueeze(review, 1) #reshape review, which originally is in seq len, input size shape
18         #to have a first axis with size 1.
19
20         out, h = model(review, h)
21         target = torch.argmax(sentiment)
22         #print(out.shape, h.shape)
23         #compute loss on output from GRU
24         loss = criterion(out[-1,:], target)
25         running_loss += loss.item()
26         loss.backward()
27         optimizer.step()
28         if n % 200 == 199:
29             avg_loss = running_loss / float(200)
30             training_loss_tally_nnGRU.append(avg_loss)
31             current_time = time.perf_counter()
32             time_elapsed = current_time-start_time
33             print("[epoch:%d iter:%d elapsed_time:%d secs]      loss: %.5f" % (epoch+1,n+1, time_elapsed,avg_loss))
34             running_loss = 0.0
35
36
```

```

1  ### Prediction / Eval Code for nn.GRU model
2  correct_pos_sentiment = 0.0
3  total_pos_sentiment = 0.0
4  correct_neg_sentiment = 0.0
5  total_neg_sentiment = 0.0
6  sentiment_pred = []
7  sentiment_label = []
8  mapping = { 0 : 'negative', 1: 'positive'}
9  with torch.no_grad():
10     for n, data in enumerate(text_cl.dataserver_test):
11         #print('n is', n)
12         sentiment = data['sentiment']
13         review = data['review']
14         h = model.init_hidden()
15         review = torch.unsqueeze(review, 1)
16         out, h = model(review, h) #goes through one sequence of embeddings (i.e one review)
17
18         GT = torch.argmax(sentiment).item()
19         sentiment_label.append(mapping[GT])
20         pred = torch.argmax(out[-1, :]).item()
21         sentiment_pred.append(mapping[pred])
22
23         if (GT == 1): #compute positive sentiment accuracy
24             total_pos_sentiment += 1
25             if pred == GT:
26                 #print('accurate pos sentiment pred')
27                 correct_pos_sentiment += 1
28
29         if (GT == 0): #compute negative sentiment accuracy
30             total_neg_sentiment += 1
31             if pred == GT:
32                 #print('accurate neg sentiment pred')
33                 correct_neg_sentiment += 1
34
35
36     print("total acc for pos:      total acc for neg:")
37     print(correct_pos_sentiment / total_pos_sentiment, '      ', correct_neg_sentiment / total_neg_sentiment)
38     print('total accuracy:', (correct_pos_sentiment + correct_neg_sentiment) / len(sentiment_label))
39
40 from sklearn.metrics import confusion_matrix
41 y_true = sentiment_label
42 y_pred = sentiment_pred
43 print(len(y_true), len(y_pred))
44 confusion_matrix=confusion_matrix(y_true, y_pred, labels = [ "positive", "negative"])
45 disp = ConfusionMatrixDisplay(confusion_matrix, display_labels = [ "positive", "negative"])
46 disp.plot()
47 disp.ax_.set_title("Confusion Matrix for Sentiment Predictor using nn.GRU - bidirectional=False")
48 plt.show()

```

Network, Training, and Evaluation code for Bidirectional GRU:

```

1  ### Train using PyTorch's bidirectional GRU:
2
3  class GRUNetPyTorch_BD(nn.Module):
4      def __init__(self, input_size, hidden_size, output_size, num_layers):
5          """
6              -- input_size is the size of the tensor for each word in a sequence of words. If you word2vec
7                  embedding, the value of this variable will always be equal to 300.
8              -- hidden_size is the size of the hidden state in the RNN
9              -- output_size is the size of output of the RNN. For binary classification of
10                  input text, output_size is 2.
11              -- num_layers creates a stack of GRUs
12          """
13          super(GRUNetPyTorch_BD, self).__init__()
14          self.input_size = input_size
15          self.hidden_size = hidden_size
16          self.num_layers = num_layers
17          self.output_size = output_size
18          self.gru = nn.GRU(self.input_size, self.hidden_size, self.num_layers, bidirectional = True)
19          self.fc = nn.Linear(self.hidden_size * 2, self.output_size)
20          self.relu = nn.ReLU()
21          self.logsoftmax = nn.LogSoftmax(dim=1)
22
23      def forward(self, x, h):
24          #print("in input: x shp and then h shape:", x.shape, h.shape)
25          out, h = self.gru(x, h)
26          #print("after running through self.gru:", out.shape, h.shape)
27          #print("out[:, -1].shape", out[:, -1].shape)
28          out = self.fc(self.relu(out[:, -1]))
29
30          #print("after relu applied to out[:, -1] and a fc(hidden size --> output size):", out.shape)
31          out = self.logsoftmax(out)
32          #print("output after logsoftmax:", out.shape)
33          return out, h
34
35      def init_hidden(self):
36          weight = next(self.parameters()).data
37          #           num_layers   batch_size   hidden_size
38          # I think that the first argument in this case refers to num_dirs
39          hidden = weight.new(  self.num_layers * 2,           1,           self.hidden_size    ).zero_()
40          return hidden
41

```

April 16, 2023

```
50 model_bidirectional = GRUNetPyTorch_BD(input_size = 300, hidden_size = 100, output_size = 2, num_layers=2)
51 epochs = 4
52 criterion = nn.NLLLoss()
53 optimizer = optim.Adam(model_bidirectional.parameters(), lr=dls.learning_rate)
54 running_loss = 0.0
55 training_loss_tally_nnGRU_bidirectional = []
56 start_time = time.perf_counter()
57
58 for epoch in range(epochs):
59     for n, data in enumerate(text_cl.dataserver_train):
60         review = data['review']
61         sentiment = data['sentiment']
62
63         #input to a GRU instance will have the shape num of embeddings, batch size, embedding size
64         review = torch.unsqueeze(review, 1)
65         #print(review.shape)
66
67         #init hidden state. When using bidirectional mode, we need to have hidden be the shape of num dirs, batch s
68         hidden = model_bidirectional.init_hidden()
69         #print(hidden.shape)
70
71         out, h = model_bidirectional(review, hidden)
72         #print("out shape:", out.shape)
73         #print("h shape", h.shape)
74         #print(out[:, -1].shape)
75
76         target = torch.argmax(sentiment)
77
78         #compute loss:
79         loss = criterion(out[:, -1], target)
80         running_loss += loss.item()
81
82         #compute gradients, backprop
83         loss.backward()
84
85         #gradient step:
86         optimizer.step()
87
88         if (n % 200) == 199:
89             avg_loss = running_loss / float(200)
90             training_loss_tally_nnGRU_bidirectional.append(avg_loss)
91             current_time = time.perf_counter()
92             time_elapsed = current_time - start_time
93             print("[epoch:%d iter:%d elapsed_time:%d secs] loss: %.5f" % (epoch+1, n+1, time_elapsed, avg_loss))
94             running_loss = 0.0
95
```

```

1  ### Prediction / Eval Code for nn.GRU model
2  correct_pos_sentiment = 0.0
3  total_pos_sentiment = 0.0
4  correct_neg_sentiment = 0.0
5  total_neg_sentiment = 0.0
6  sentiment_pred = []
7  sentiment_label = []
8  mapping = { 0 : 'negative', 1: 'positive'}
9  with torch.no_grad():
10     for n, data in enumerate(text_cl.dataserver_test):
11
12         #print('n is', n)
13         sentiment = data['sentiment']
14         review = data['review']
15         h = model_bidirectional.init_hidden()
16         review = torch.unsqueeze(review, 1)
17         out, h = model_bidirectional(review, h) #goes through one sequence of embeddings (i.e one review)
18
19         GT = torch.argmax(sentiment).item()
20         sentiment_label.append(mapping[GT])
21         pred = torch.argmax(out[-1, :]).item()
22         sentiment_pred.append(mapping[pred])
23         if (GT == 1): #compute positive sentiment accuracy
24             total_pos_sentiment += 1
25             if pred == GT:
26                 #print('accurate pos sentiment pred')
27                 correct_pos_sentiment += 1
28
29         if (GT == 0): #compute negative sentiment accuracy
30             total_neg_sentiment += 1
31             if pred == GT:
32                 #print('accurate neg sentiment pred')
33                 correct_neg_sentiment += 1
34
35
36         print("total acc for pos:      total acc for neg:")
37         print(correct_pos_sentiment / total_pos_sentiment, ' ', correct_neg_sentiment / total_neg_sentiment)
38         print('total accuracy:', (correct_pos_sentiment + correct_neg_sentiment) / len(sentiment_label))
39
40 from sklearn.metrics import confusion_matrix
41 y_true = sentiment_label
42 y_pred = sentiment_pred
43 print(len(y_true), len(y_pred))
44 confusion_matrix=confusion_matrix(y_true, y_pred, labels = [ "positive", "negative"])
45 disp = ConfusionMatrixDisplay(confusion_matrix, display_labels = [ "positive", "negative"])
46 disp.plot()
47 disp.ax_.set_title("Confusion Matrix for Sentiment Predictor using Bidirectional GRU")
48 plt.show()

```